

MP4

April 19, 2023

1 Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on a dataset of cat face images.

```
[1]: import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2
```

```
[2]: from gan.train import train
```

```
[3]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

2 GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

2.0.1 GAN loss

TODO: Implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log (1 - D(G(z)))]$$

Note that these are negated from the equations presented earlier as we will be *minimizing* these losses.

HINTS: You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score $s \in \mathbb{R}$ and a label $y \in \{0, 1\}$, the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of $\log D(G(z))$, $\log D(x)$ and $\log(1 - D(G(z)))$, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

```
[4]: from gan.losses import discriminator_loss, generator_loss
```

2.0.2 Least Squares GAN loss

TODO: Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

HINTS: Instead of computing the expectation, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing. When plugging in for $D(x)$ and $D(G(z))$ use the direct output from the discriminator (`scores_real` and `scores_fake`).

```
[5]: from gan.losses import ls_discriminator_loss, ls_generator_loss
```

3 GAN model architecture

TODO: Implement the Discriminator and Generator networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

Discriminator:

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`
- convolutional layer with `in_channels=128, out_channels=256, kernel=4, stride=2`
- batch norm

- convolutional layer with in_channels=256, out_channels=512, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=512, out_channels=1024, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=1024, out_channels=1, kernel=4, stride=1

Use padding = 1 (not 0) for all the convolutional layers.

Instead of Relu we LeakyReLU throughout the discriminator (we use a negative slope value of 0.2). You can use simply use relu as well.

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

Generator:

Note: In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1
- batch norm
- transpose convolution with in_channels=1024, out_channels=512, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=512, out_channels=256, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=256, out_channels=128, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=128, out_channels=3, kernel=4, stride=2

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a 3x64x64 tensor for each sample (equal dimensions to the images from the dataset).

```
[6]: from gan.models import Discriminator, Generator
```

4 Data loading

The cat images we provide are RGB images with a resolution of 64x64. In order to prevent our discriminator from overfitting, we will need to perform some data augmentation.

TODO: Implement data augmentation by adding new transforms to the cell below. At the minimum, you should have a RandomCrop and a ColorJitter, but we encourage you to experiment with different augmentations to see how the performance of the GAN changes. See <https://pytorch.org/vision/stable/transforms.html>.

```
[7]: #!/sh download_cat.sh
```

```
[8]: batch_size = 32
imsize = 64
```

```

cat_root = './cats'

cat_train = ImageFolder(root=cat_root, transform=transforms.Compose([
    transforms.ToTensor(),

    # Example use of RandomCrop:
    transforms.Resize(int(1.15 * imsize)),
    transforms.RandomCrop(imsize),
    transforms.ColorJitter(brightness=0.1, contrast=0.3, saturation=0.4, hue=-0.05), #Color Jitter
    #transforms.RandomAffine(degrees=15, translate=(0.1, 0.1), scale=(0.9, 1.1), shear=10),
    transforms.GaussianBlur(kernel_size=3),
])))

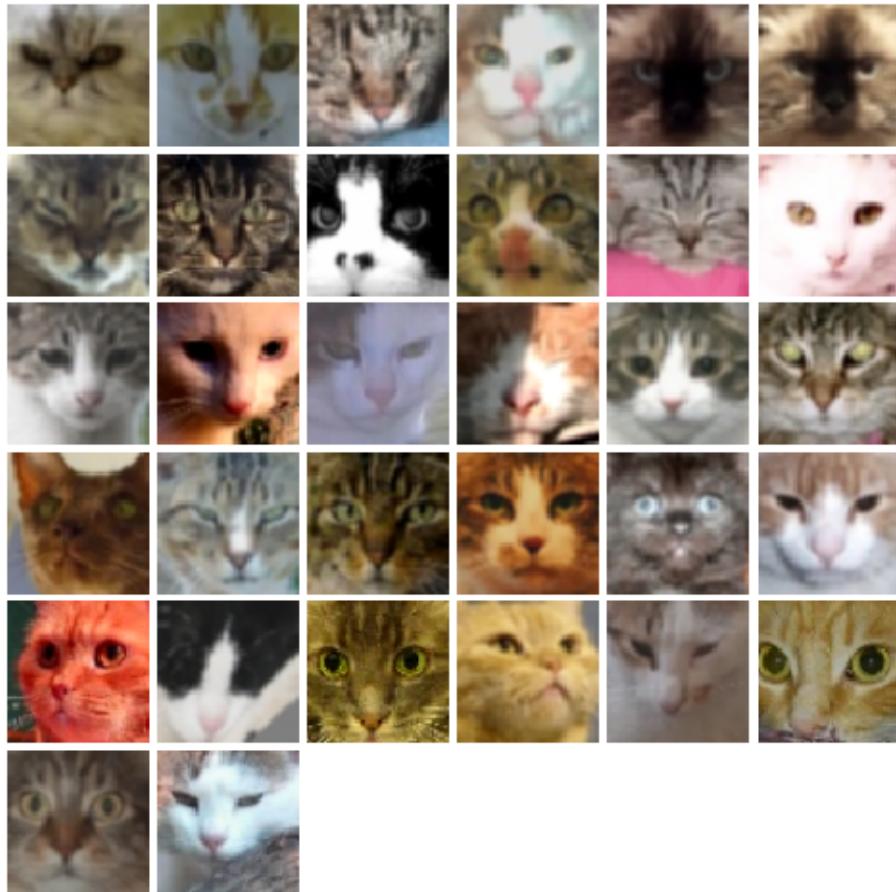
cat_loader_train = DataLoader(cat_train, batch_size=batch_size, drop_last=True)

```

4.0.1 Visualize dataset

```
[9]: from gan.utils import show_images
imgs = next(cat_loader_train.__iter__())[0].numpy().squeeze()
show_images(imgs, color=True)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



5 Training

TODO: Fill in the training loop in `gan/train.py`.

```
[10]: # NOISE_DIM = 100
      # NUM_EPOCHS = 50
      # learning_rate = 0.001
```

5.0.1 Train GAN

```
[11]: # D = Discriminator().to(device)
      # G = Generator(noise_dim=NOISE_DIM).to(device)
```

```
[12]: # D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.
      ↵5, 0.999))
      # G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.
      ↵5, 0.999))
```

```
[13]: # # original gan
# train(D, G, D_optimizer, G_optimizer, discriminator_loss,
#        generator_loss, num_epochs=NUM_EPOCHS, show_every=1000,
#        batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

5.0.2 Train LS-GAN

```
[14]: # D = Discriminator().to(device)
# G = Generator(noise_dim=NOISE_DIM).to(device)
```

```
[15]: # D_optimizer = torch.optim.Adam(D.parameters(), lr=1e-3, betas = (0.5, 0.999))
# G_optimizer = torch.optim.Adam(G.parameters(), lr=1e-3, betas = (0.5, 0.999))
```

```
[16]: # # ls-gan
# train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
#        ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=1000,
#        batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

Extra Credit 1: Train Wasserstein GAN

```
[17]: # D = Discriminator().to(device)
# G = Generator(noise_dim=NOISE_DIM).to(device)
```

```
[18]: # D_optimizer = torch.optim.Adam(D.parameters(), lr=5e-4, betas = (0.5, 0.999))
# G_optimizer = torch.optim.Adam(G.parameters(), lr=5e-4, betas = (0.5, 0.999))
```

```
[19]: # # Wasserstein-gan
# from gan.losses import wass_discriminator_loss, wass_generator_loss
# train(D, G, D_optimizer, G_optimizer, wass_discriminator_loss,
#        wass_generator_loss, num_epochs=40, show_every=1000,
#        batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

Extra Credit 2: Different Dataset - Celeb Faces

```
[20]: # batch_size = 128
# imsize = 64
# faces_root = './celeb'
# faces_train = ImageFolder(root=faces_root, transform=transforms.Compose([
#     transforms.ToTensor(),
#
#     # Example use of RandomCrop:
#     transforms.Resize(int(1.1 * imsize)),
#     transforms.RandomCrop(imsize),
#     transforms.ColorJitter(brightness=0.1, contrast=0.3, saturation=0.4, hue=0.05),
#     transforms.GaussianBlur(kernel_size=5)
# ]))
```

```
# faces_loader_train = DataLoader(faces_train, batch_size=batch_size, drop_last=True)
```

Visualize this Dataset

```
[21]: # from gan.utils import show_images
# imgs = next(faces_loader_train.__iter__())[0].numpy().squeeze()
# show_images(imgs, color=True)
```

Train with LS-GAN with Faces Dataset

```
[22]: # D = Discriminator().to(device)
# G = Generator(noise_dim=NOISE_DIM).to(device)
```

```
[23]: # D_optimizer = torch.optim.Adam(D.parameters(), lr=1e-4, betas = (0.5, 0.999))
# G_optimizer = torch.optim.Adam(G.parameters(), lr=1e-4, betas = (0.5, 0.999))
```

```
[24]: # # ls-gan
# train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
#        ls_generator_loss, num_epochs=30, show_every=1000,
#        batch_size=batch_size, train_loader=faces_loader_train, device=device)
```

Extra Credit 3: Training with LSGAN with Spectral Normalization on Cat Faces Dataset

```
[25]: NOISE_DIM = 100
NUM_EPOCHS = 30
```

```
[26]: D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)
```

```
[27]: D_optimizer = torch.optim.Adam(D.parameters(), lr=1e-4, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=1e-4, betas = (0.5, 0.999))
```

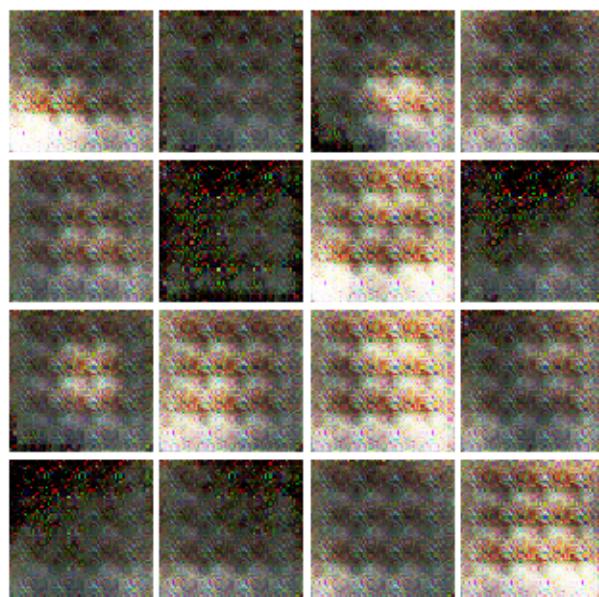
```
[28]: # ls-gan
train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
      ls_generator_loss, num_epochs=50, show_every=500,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

EPOCH: 1
Iter: 0, D: 0.5701, G:0.7692



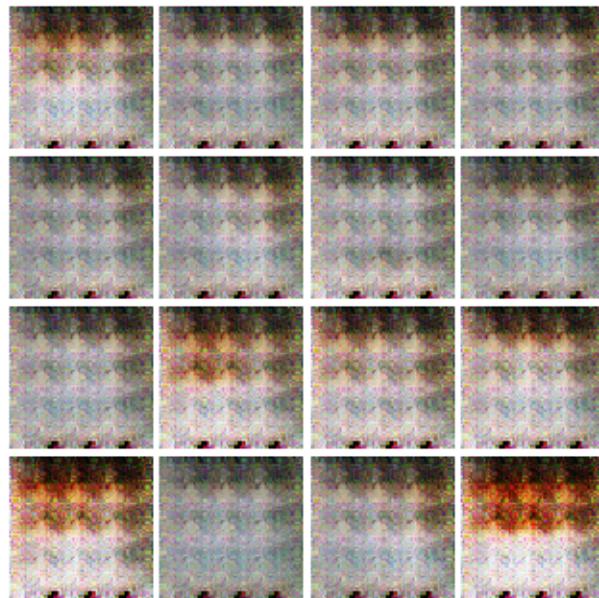
EPOCH: 2

Iter: 500, D: 0.03154, G: 0.5507



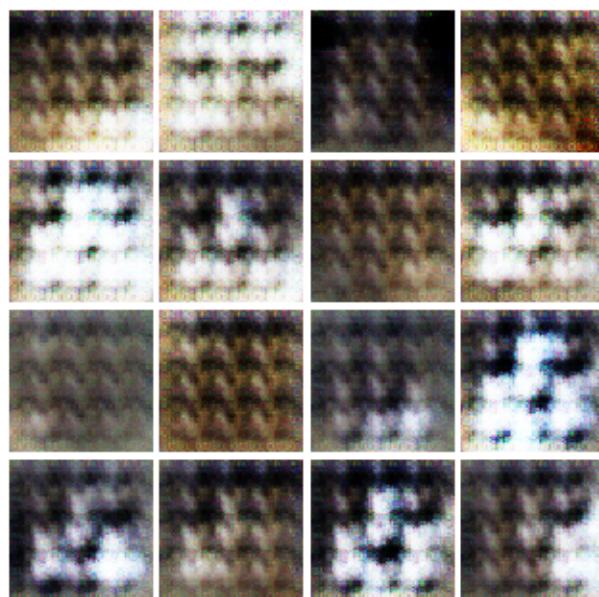
EPOCH: 3

Iter: 1000, D: 0.02893, G: 0.5888



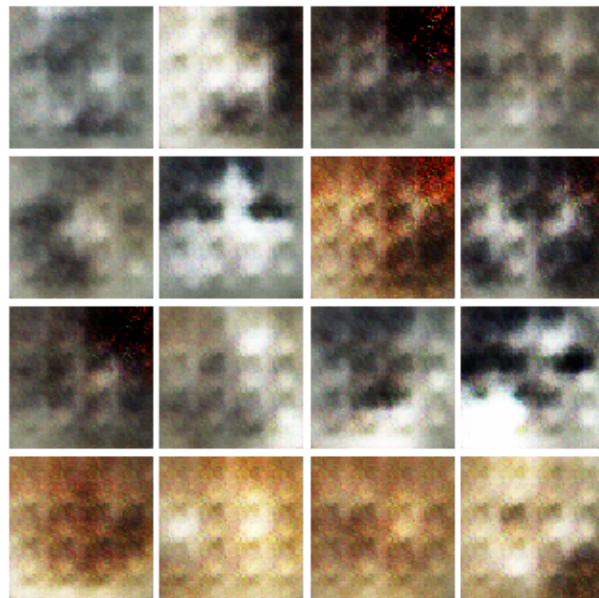
EPOCH: 4

Iter: 1500, D: 0.01857, G: 0.5524



EPOCH: 5

Iter: 2000, D: 0.0125, G: 0.5872



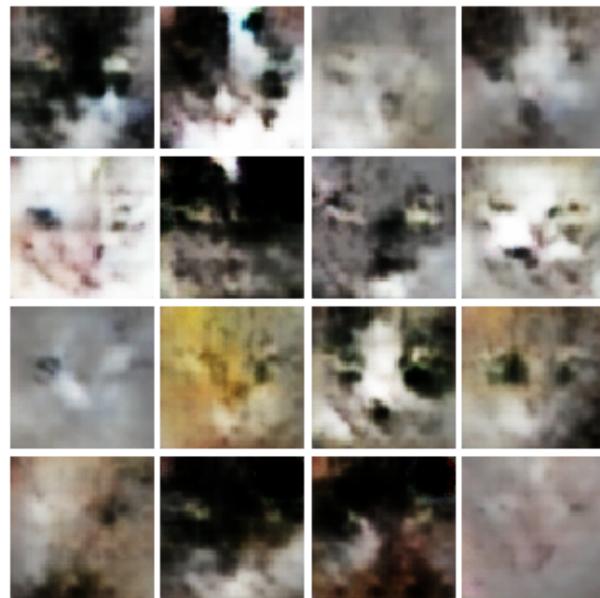
EPOCH: 6

Iter: 2500, D: 0.03684, G: 0.5476



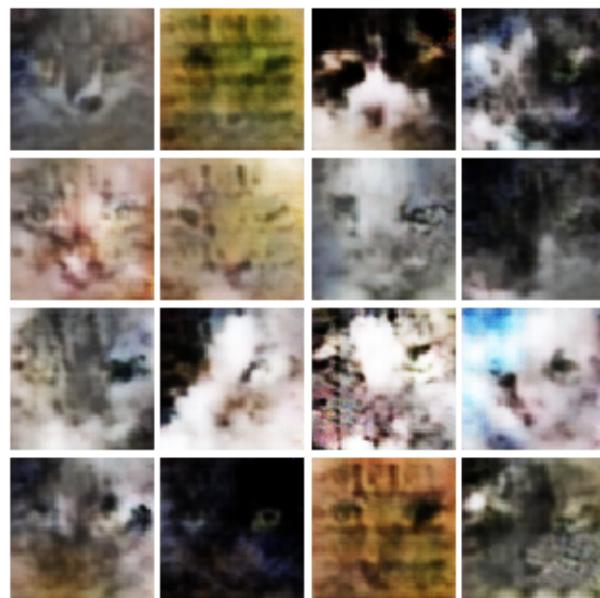
EPOCH: 7

Iter: 3000, D: 0.04791, G: 0.5042



EPOCH: 8

Iter: 3500, D: 0.04704, G: 0.3508



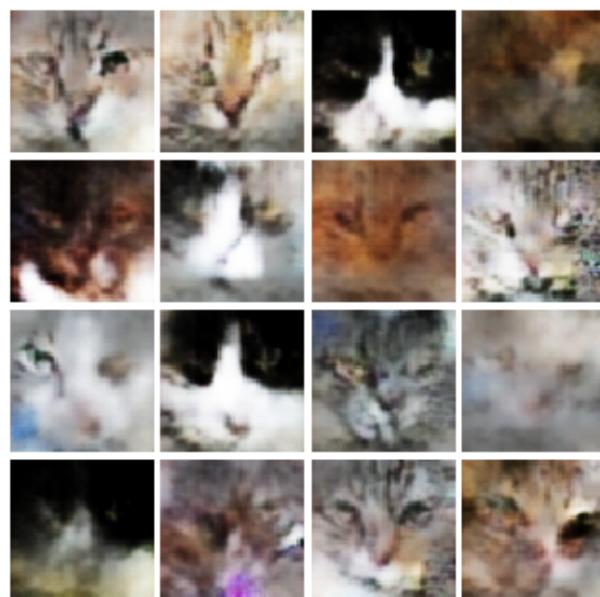
EPOCH: 9

Iter: 4000, D: 0.03862, G: 0.5534



EPOCH: 10

Iter: 4500, D: 0.03867, G: 0.676



EPOCH: 11

Iter: 5000, D: 0.0209, G: 0.5145



EPOCH: 12

Iter: 5500, D: 0.02054, G: 0.5657



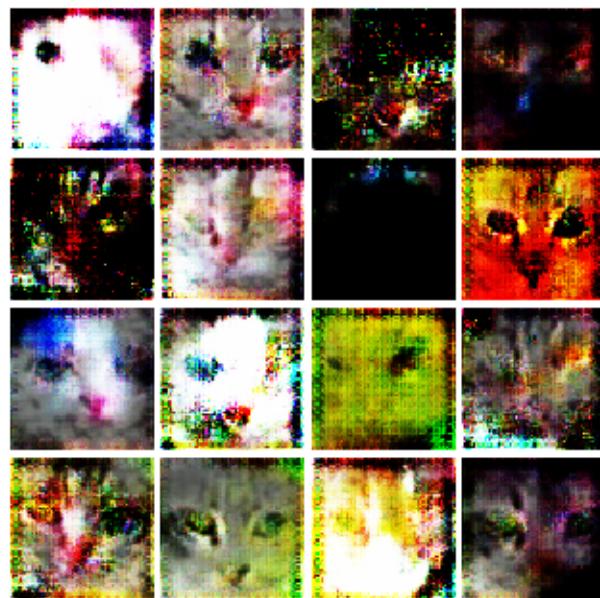
EPOCH: 13

Iter: 6000, D: 0.01758, G: 0.5123



EPOCH: 14

Iter: 6500, D: 0.03502, G: 0.5786



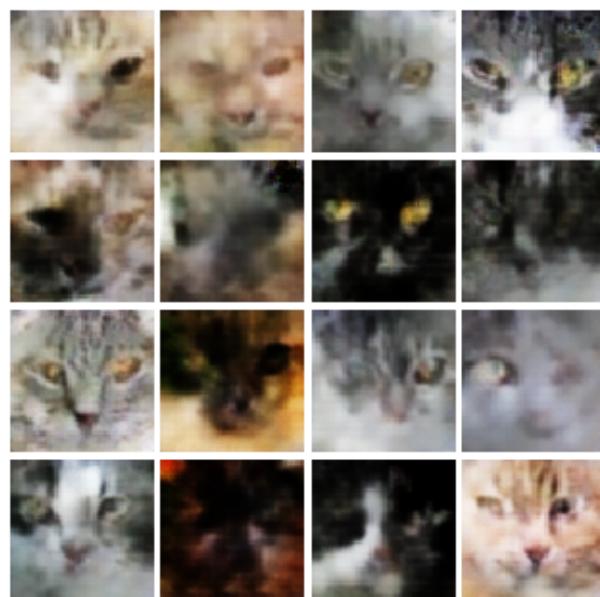
EPOCH: 15

Iter: 7000, D: 0.0157, G: 0.4764



EPOCH: 16

Iter: 7500, D: 0.02392, G: 0.5337



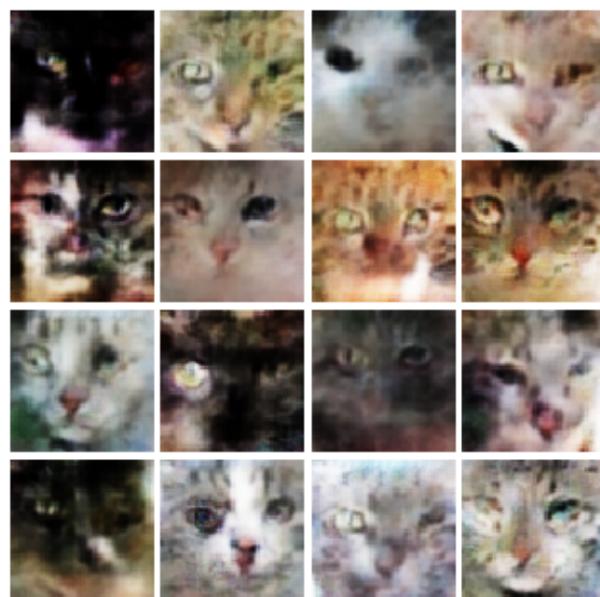
EPOCH: 17

Iter: 8000, D: 0.07737, G: 0.6442



EPOCH: 18

Iter: 8500, D: 0.04656, G:0.5321



EPOCH: 19

Iter: 9000, D: 0.06752, G:0.3259



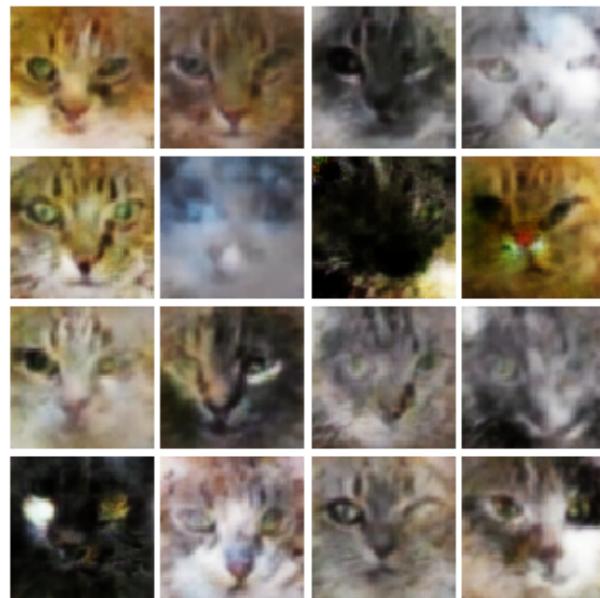
EPOCH: 20

Iter: 9500, D: 0.01408, G: 0.531



EPOCH: 21

Iter: 10000, D: 0.01728, G: 0.5351



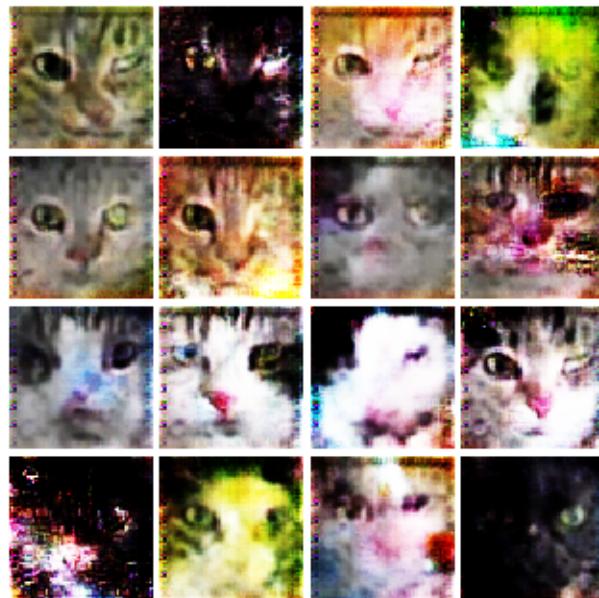
EPOCH: 22

Iter: 10500, D: 0.02749, G: 0.4439



EPOCH: 23

Iter: 11000, D: 0.01676, G: 0.5135



EPOCH: 24

Iter: 11500, D: 0.003558, G:0.53



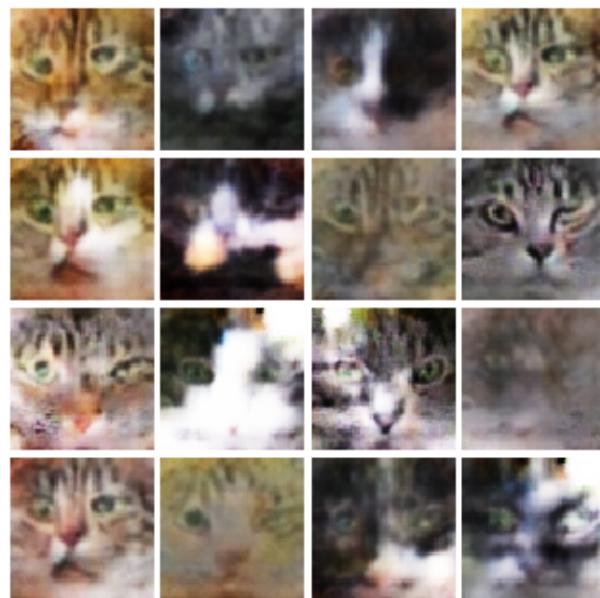
EPOCH: 25

Iter: 12000, D: 0.002614, G:0.501



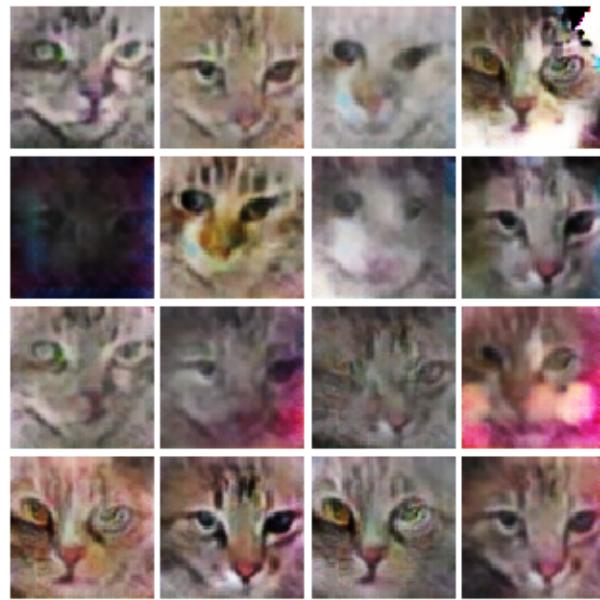
EPOCH: 26

Iter: 12500, D: 0.005064, G:0.5133



EPOCH: 27

Iter: 13000, D: 0.002759, G:0.5314



EPOCH: 28

Iter: 13500, D: 0.001112, G:0.5204



EPOCH: 29

Iter: 14000, D: 0.02201, G:0.535



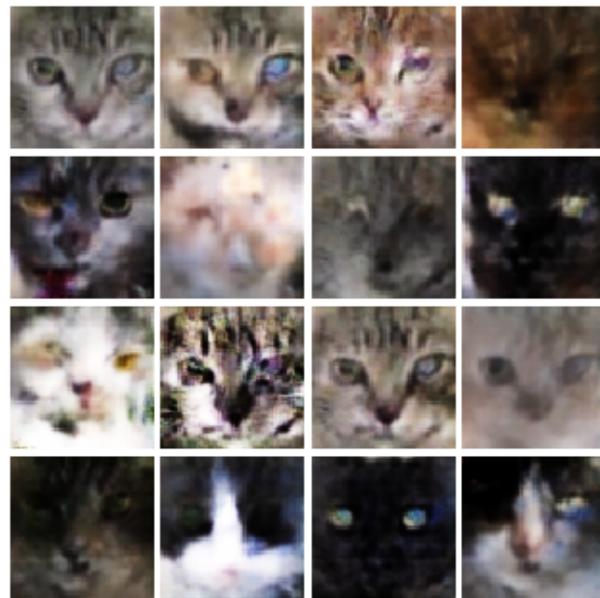
EPOCH: 30

Iter: 14500, D: 0.006943, G:0.5638



EPOCH: 31

Iter: 15000, D: 0.007305, G:0.4898



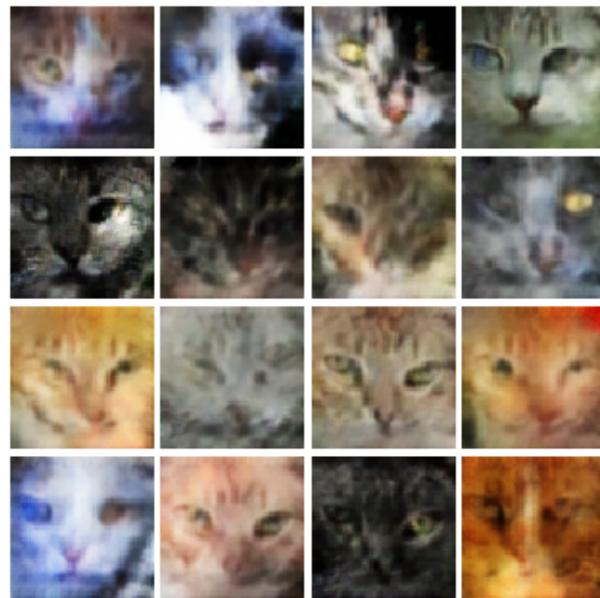
EPOCH: 32

Iter: 15500, D: 0.03835, G: 0.6011



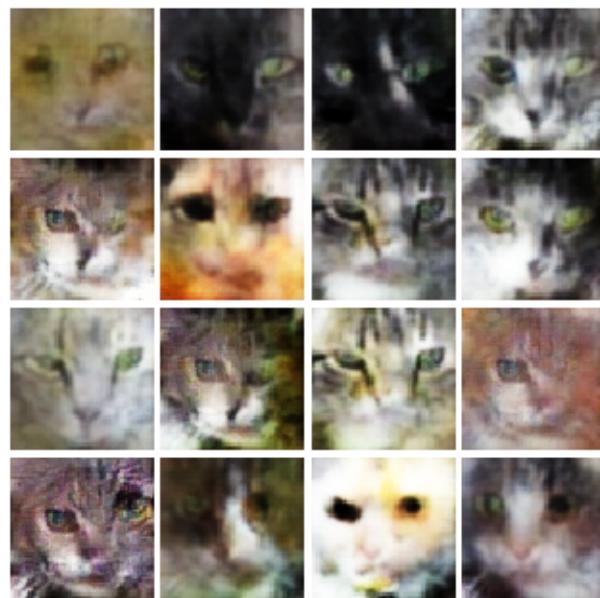
EPOCH: 33

Iter: 16000, D: 0.01494, G: 0.4709



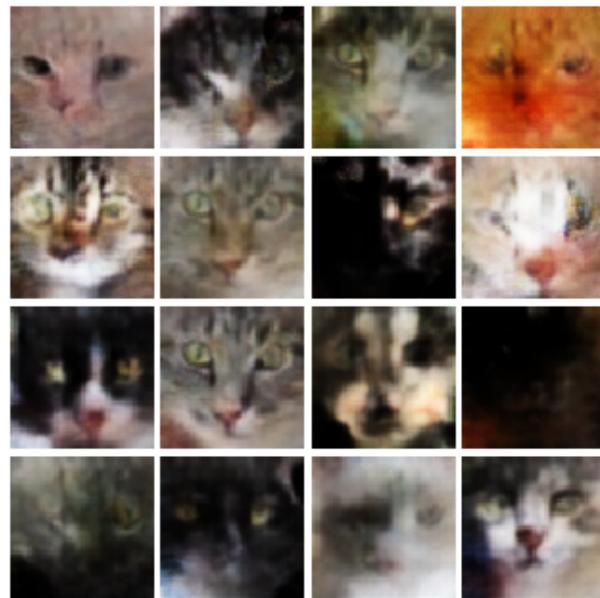
EPOCH: 34

Iter: 16500, D: 0.008395, G:0.5655



EPOCH: 35

Iter: 17000, D: 0.01234, G:0.5688



EPOCH: 36

Iter: 17500, D: 0.03543, G: 0.5059



EPOCH: 37

Iter: 18000, D: 0.01065, G: 0.5584



EPOCH: 38

Iter: 18500, D: 0.04469, G:0.5948



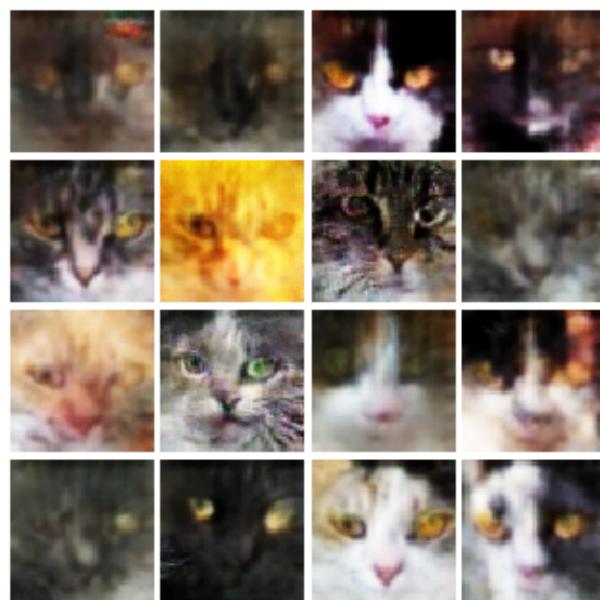
EPOCH: 39

Iter: 19000, D: 0.01697, G:0.449



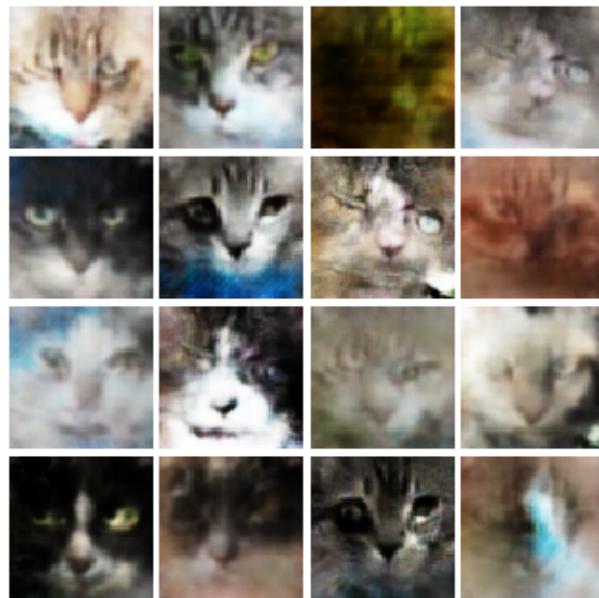
EPOCH: 40

Iter: 19500, D: 0.01192, G:0.5329



EPOCH: 41

Iter: 20000, D: 0.01197, G:0.4984



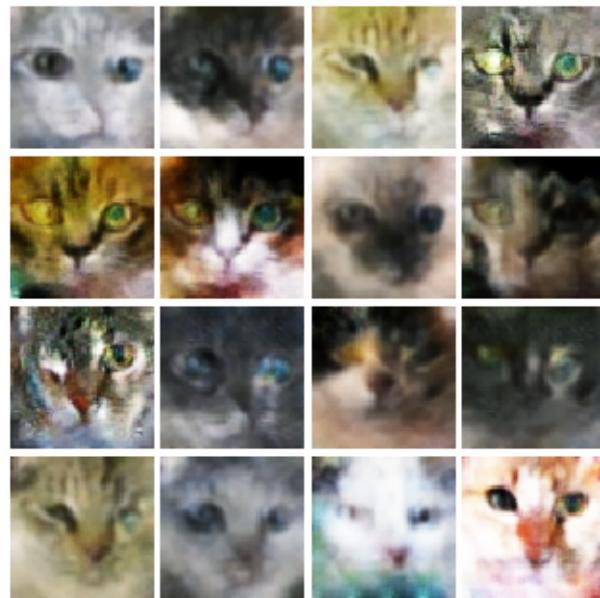
EPOCH: 42

Iter: 20500, D: 0.01079, G:0.5165



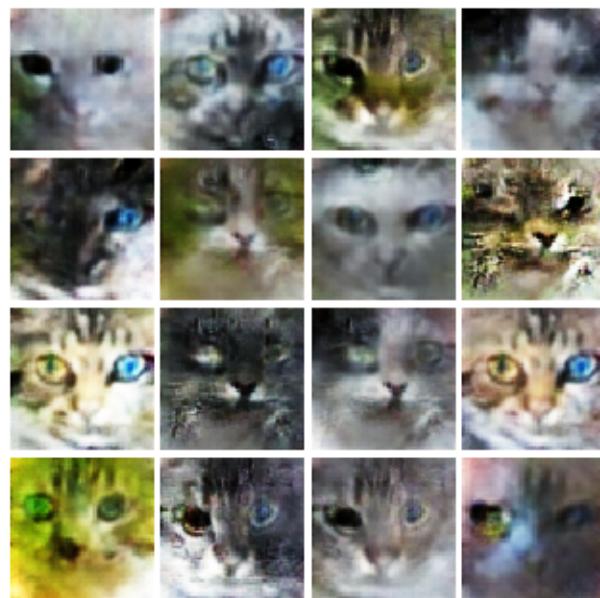
EPOCH: 43

Iter: 21000, D: 0.0178, G:0.577



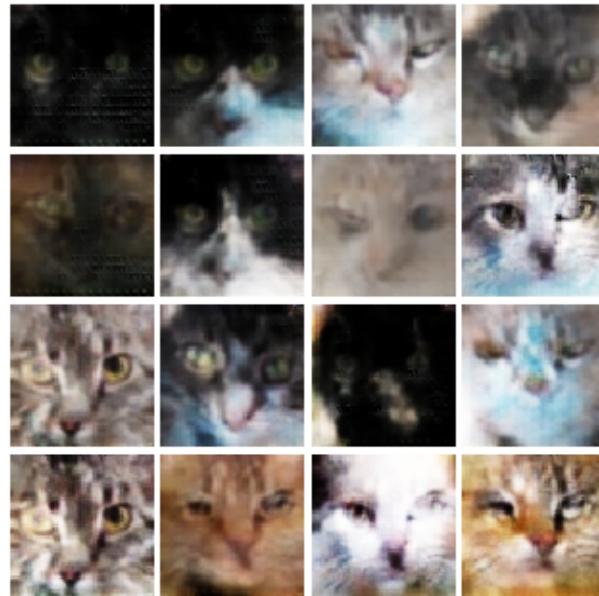
EPOCH: 44

Iter: 21500, D: 0.006447, G:0.5684



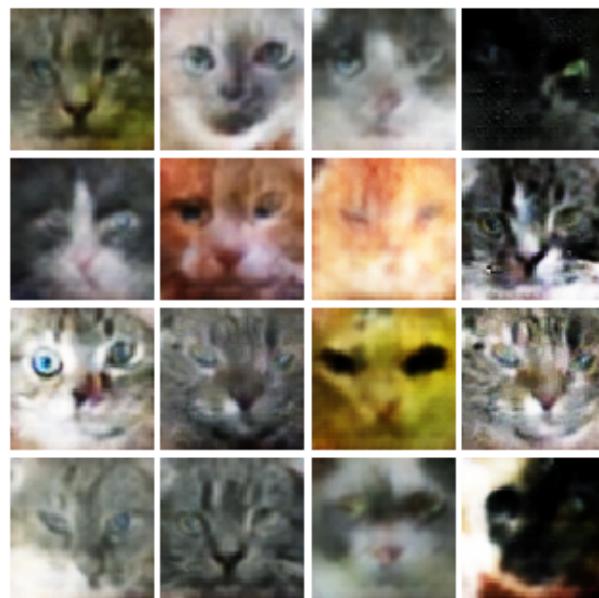
EPOCH: 45

Iter: 22000, D: 0.01737, G:0.604



EPOCH: 46

Iter: 22500, D: 0.007623, G:0.5405



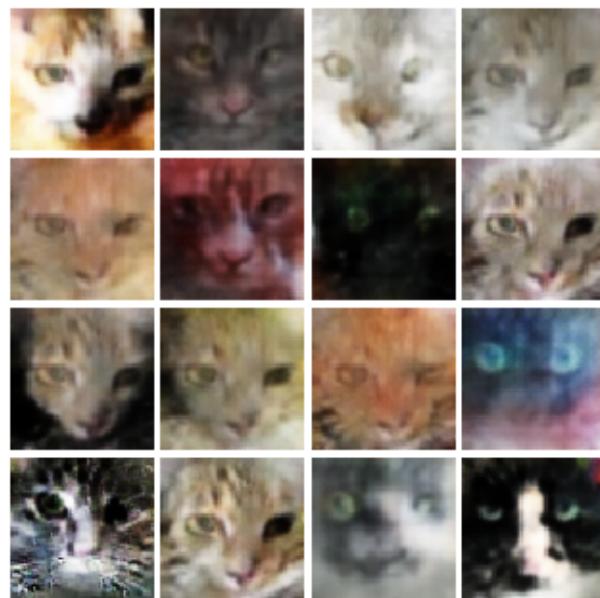
EPOCH: 47

Iter: 23000, D: 0.01481, G:0.5264



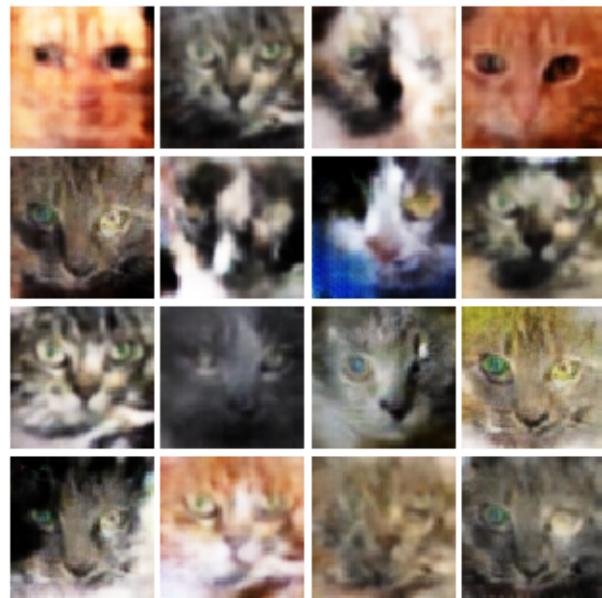
EPOCH: 48

Iter: 23500, D: 0.004204, G:0.5395



EPOCH: 49

Iter: 24000, D: 0.006177, G:0.5456



EPOCH: 50

Iter: 24500, D: 0.01253, G: 0.5796

